

Alex Welcing

alexwelcing@gmail.com · welc.ing · github.com/alexwelcing · linkedin.com/in/alexwelcing

NOTION · FORWARD DEPLOYED ENGINEER, GTM

Hiring Team at Notion

I am applying for the Forward Deployed Engineer, GTM role (San Francisco, CA).

My pitch for this role is simple: I build AI products as working systems: multi-agent orchestration across multiple LLM providers, a type-safe AI API execution workspace with evals and live introspection, and enterprise deployments where reliability is non-negotiable.

WHY I AM A STRONG FIT:

- Production agents, not prototypes: Built a multi-agent orchestration platform end-to-end across multiple LLM providers, with git-backed persistence and eval loops that improved accuracy across three internal APIs.
- Platform judgment from billions-scale systems: Re-platformed identity and subscription infrastructure at ALM serving billions of monthly requests — I know what production-grade means for AI surfaces.
- Customer-facing technical credibility: Enterprise SSO onboarding for 150+ firms and 3D-commerce deployments for Alo, Moncler, and Ralph Lauren — I embed with customers and feed patterns back to product.

A FEW SHIPPED REFERENCE POINTS:

- Multi-agent orchestration platform built end-to-end across multiple LLM providers
- Type-safe AI API execution workspace shipped at ALM with live introspection and evals
- MCP servers, sub-agents, and skills shipped as independent systems

Packet (résumé, cover letter, and a 30-second role pitch): welc.ing/r/pkt-fi4n8s

I would welcome the opportunity to discuss how I can contribute quickly, operate with high autonomy, and ship for this team.

Best regards,
Alex Welcing
alexwelcing@gmail.com
welc.ing/r/pkt-fi4n8s