
Alex Welcing

alexwelcing@gmail.com · welc.ing · github.com/alexwelcing · linkedin.com/in/alexwelcing

EXA · FORWARD DEPLOYED ENGINEER

Hiring Team at Exa

I am applying for the Forward Deployed Engineer role (San Francisco, CA).

My pitch for this role is simple: Architect PM who builds: lupi.live and lupine.science shipped solo as a materials-science AI stack, a marketing-automation OS deployed on GCP, and enterprise platforms rebuilt at billions-scale — 12+ years of end-to-end ownership.

WHY I AM A STRONG FIT:

- Zero-to-one range, repeatedly proven: Shipped lupi.live (R3F/WebGL molecule viewer), lupine.science (MLIP research layer), and a GCP marketing-automation OS — each solo, from concept to production.
- Enterprise-scale judgment when it matters: Re-platformed identity infrastructure at ALM serving billions of monthly requests — startup pace with big-system discipline.
- AI-native building style: Multi-agent orchestration across LLM providers, evals, document AI, and semantic search — AI is my default toolkit, not an add-on.

A FEW SHIPPED REFERENCE POINTS:

- 12+ years shipping product, code, and judgment end-to-end
- lupi.live and lupine.science: materials-science AI stack shipped solo
- 3D web, semantic search, and scientific computing in production

Packet (résumé, cover letter, and a 30-second role pitch): welc.ing/r/pkt-1hjbgw8

I would welcome the opportunity to discuss how I can contribute quickly, operate with high autonomy, and ship for this team.

Best regards,
Alex Welcing
alexwelcing@gmail.com
welc.ing/r/pkt-1hjbgw8